

Problem Statement

You are hired by one of the leading news channel CNBE who wants to analyze recent elections. This survey was conducted on 1525 voters with 9 variables. You have to build a model, to predict which party a voter will vote for on the basis of the given information, to create an exit poll that will help in predicting overall win and seats covered by a particular party. Our Dataset has following variables:

Data Dictionary

1. vote: Party choice: Conservative or Labour
2. age: in years
3. economic.cond.national: Assessment of current national economic conditions, 1 to 5.
4. economic.cond.household: Assessment of current household economic conditions, 1 to 5.
5. Blair: Assessment of the Labour leader, 1 to 5.
6. Hague: Assessment of the Conservative leader, 1 to 5.
7. Europe: an 11-point scale that measures respondents' attitudes toward European integration. High scores represent 'Eurosceptic' sentiment.
8. political.knowledge: Knowledge of parties' positions on European integration, 0 to 3.
9. gender: female or male.

1. Read the dataset. Do the descriptive statistics and do null value condition check. Write an inference on it. (5 Marks)

Upload Required Libraries

```
In [965]: import numpy as np
import pandas as pd
import os

import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.style
plt.style.use('classic')

import warnings
warnings.filterwarnings("ignore")
```

Importing data

```
In [966]: ## Load the csv file available in the working or specified directory

df = pd.read_excel("Election_Data.xlsx")
```

EDA

```
In [967]: # Check top few records to get a feel of the data structure

df.head()
```

```
Out[967]:
```

	Unnamed: 0	vote	age	economic.cond.national	economic.cond.household	Blair	Hague	I
0	1	Labour	43	3	3	4	1	
1	2	Labour	36	4	4	4	4	
2	3	Labour	35	4	4	5	2	
3	4	Labour	24	4	2	2	1	
4	5	Labour	41	2	2	1	1	

Shape

```
In [968]: print("No of rows: ",df.shape[0], "\n""No of columns: ",df.shape[1])
```

```
No of rows: 1525
No of columns: 10
```

Data type of data features

```
In [969]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1525 entries, 0 to 1524
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            1525 non-null   int64
1   vote                                  1525 non-null   object
2   age                                   1525 non-null   int64
3   economic.cond.national                1525 non-null   int64
4   economic.cond.household               1525 non-null   int64
5   Blair                                 1525 non-null   int64
6   Hague                                  1525 non-null   int64
7   Europe                                1525 non-null   int64
8   political.knowledge                   1525 non-null   int64
9   gender                                 1525 non-null   object
dtypes: int64(8), object(2)
memory usage: 119.3+ KB
```

checking for Possible columns which are categorical but are having data type "object"

```
In [970]: df['vote'].value_counts()
```

```
Out[970]: Labour          1063
Conservative      462
Name: vote, dtype: int64
```

```
In [971]: df['gender'].value_counts()
```

```
Out[971]: female      812
male          713
Name: gender, dtype: int64
```

Remove unnamed column

```
In [972]: df.drop(df.columns[0],axis=1,inplace=True)
```

In [973]:

df

Out[973]:

	vote	age	economic.cond.national	economic.cond.household	Blair	Hague	Eur
0	Labour	43	3	3	4	1	
1	Labour	36	4	4	4	4	
2	Labour	35	4	4	5	2	
3	Labour	24	4	2	2	1	
4	Labour	41	2	2	1	1	
...
1520	Conservative	67	5	3	2	4	
1521	Conservative	73	2	2	4	4	
1522	Labour	37	3	3	5	4	
1523	Conservative	61	3	3	1	4	
1524	Conservative	74	2	3	2	4	

1525 rows × 9 columns



checking for Possible columns which are categorical but are not having data type "object"

In [974]:

df['economic.cond.national'].value_counts()

Out[974]:

```
3    607
4    542
2    257
5     82
1     37
```

Name: economic.cond.national, dtype: int64

In [975]:

df['economic.cond.household'].value_counts()

Out[975]:

```
3    648
4    440
2    280
5     92
1     65
```

Name: economic.cond.household, dtype: int64

In [976]:

df['Blair'].value_counts()

Out[976]:

```
4    836
2    438
5    153
1     97
3      1
```

Name: Blair, dtype: int64

```
In [977]: df['Hague'].value_counts()
```

```
Out[977]: 2    624
          4    558
          1    233
          5     73
          3     37
          Name: Hague, dtype: int64
```

```
In [978]: df['Europe'].value_counts()
```

```
Out[978]: 11    338
          6    209
          3    129
          4    127
          5    124
          8    112
          9    111
          1    109
          10   101
          7     86
          2     79
          Name: Europe, dtype: int64
```

```
In [979]: df['political.knowledge'].value_counts()
```

```
Out[979]: 2    782
          0    455
          3    250
          1     38
          Name: political.knowledge, dtype: int64
```

Change the data types of these 6 features

***Convert political.knowledge to object type as it is having 0 which will affect in the further process ***

```
In [980]: cat=["economic.cond.national","economic.cond.household","Blair","Hague","Eur
```

```
In [981]: for i in cat:
           df[i]=df[i].astype("object")
```

In [982]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1525 entries, 0 to 1524
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   vote                                  1525 non-null   object
1   age                                   1525 non-null   int64
2   economic.cond.national               1525 non-null   object
3   economic.cond.household              1525 non-null   object
4   Blair                                1525 non-null   object
5   Hague                                 1525 non-null   object
6   Europe                                1525 non-null   object
7   political.knowledge                  1525 non-null   object
8   gender                                1525 non-null   object
dtypes: int64(1), object(8)
memory usage: 107.4+ KB
```

Making different list for categorical columns and numerical columns

```
In [983]: cat=[]
num=[]
for i in df.columns:
    if df[i].dtype=="object":
        cat.append(i)
    else:
        num.append(i)
print(cat)

print(num)
```

```
['vote', 'economic.cond.national', 'economic.cond.household', 'Blair', 'Hague', 'Europe', 'political.knowledge', 'gender']
['age']
```

In [984]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1525 entries, 0 to 1524
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   vote                                  1525 non-null   object
1   age                                   1525 non-null   int64
2   economic.cond.national               1525 non-null   object
3   economic.cond.household              1525 non-null   object
4   Blair                                1525 non-null   object
5   Hague                                 1525 non-null   object
6   Europe                                1525 non-null   object
7   political.knowledge                  1525 non-null   object
8   gender                                1525 non-null   object
dtypes: int64(1), object(8)
memory usage: 107.4+ KB
```

In [985]: `df.head()`

Out[985]:

	vote	age	economic.cond.national	economic.cond.household	Blair	Hague	Europe	poli
0	Labour	43	3	3	4	1	2	
1	Labour	36	4	4	4	4	5	
2	Labour	35	4	4	5	2	3	
3	Labour	24	4	2	2	1	4	
4	Labour	41	2	2	1	1	6	

Describe for numerical and categorical columns

In [986]: `df[num].describe().T`

Out[986]:

	count	mean	std	min	25%	50%	75%	max
age	1525.0	54.182295	15.711209	24.0	41.0	53.0	67.0	93.0

In [987]: `df[cat].describe().T`

Out[987]:

	count	unique	top	freq
vote	1525	2	Labour	1063
economic.cond.national	1525	5	3	607
economic.cond.household	1525	5	3	648
Blair	1525	5	4	836
Hague	1525	5	2	624
Europe	1525	11	11	338
political.knowledge	1525	4	2	782
gender	1525	2	female	812

In [988]: *# Are there any missing values ?*
`df.isnull().sum()`

Out[988]:

vote	0
age	0
economic.cond.national	0
economic.cond.household	0
Blair	0
Hague	0
Europe	0
political.knowledge	0
gender	0
dtype: int64	

There are no missing values

```
In [989]: ## Are there any duplicate records

# Check for duplicate data

dups = df.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))
print(df.shape)
```

```
Number of duplicate rows = 8
(1525, 9)
```

```
In [990]: df.drop_duplicates(inplace=True)
```

```
In [991]: dups = df.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))
print(df.shape)
```

```
Number of duplicate rows = 0
(1517, 9)
```

Getting unique counts of Target


```
In [992]: df.vote.value_counts(normalize=True)
```

```
Out[992]: Labour          0.69677
Conservative    0.30323
Name: vote, dtype: float64
```

2. Perform Univariate and Bivariate Analysis. Do exploratory data analysis. Check for Outliers.


Univariate Analysis

```
In [993]: fig = plt.figure(figsize = (8,6))  
ax = fig.gca()  
df.hist(ax=ax)  
plt.show()
```



```
In [994]: fig = plt.figure(figsize = (8,4))  
df.boxplot()
```

Out[994]: <AxesSubplot:>



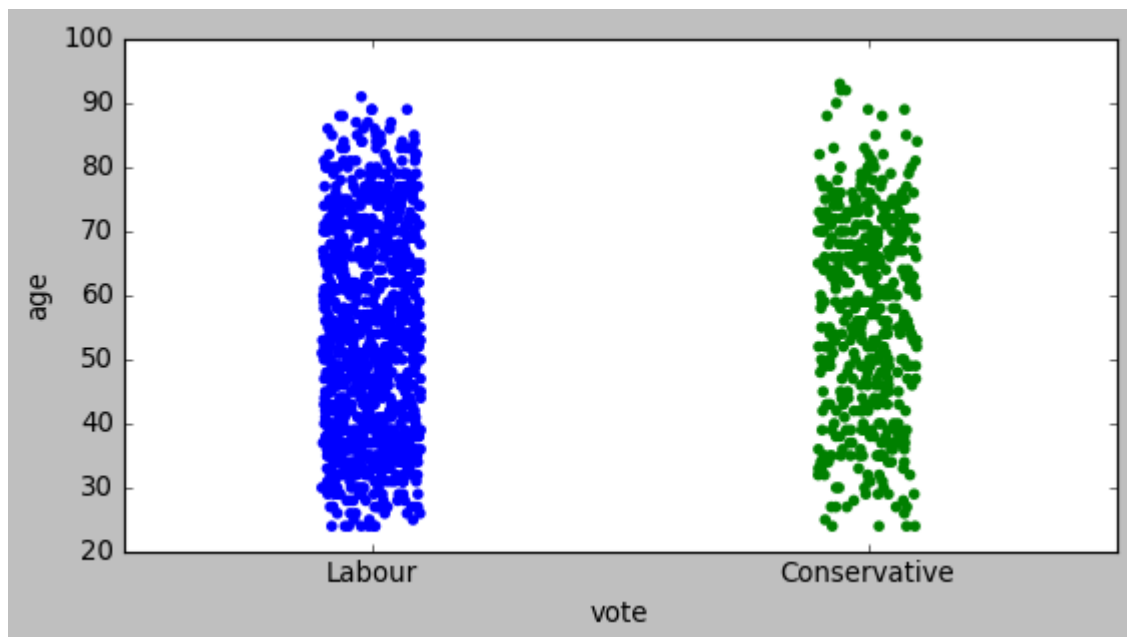
Bivariate and Multivariate Analysis

```
In [995]: print(num)
```

```
['age']
```

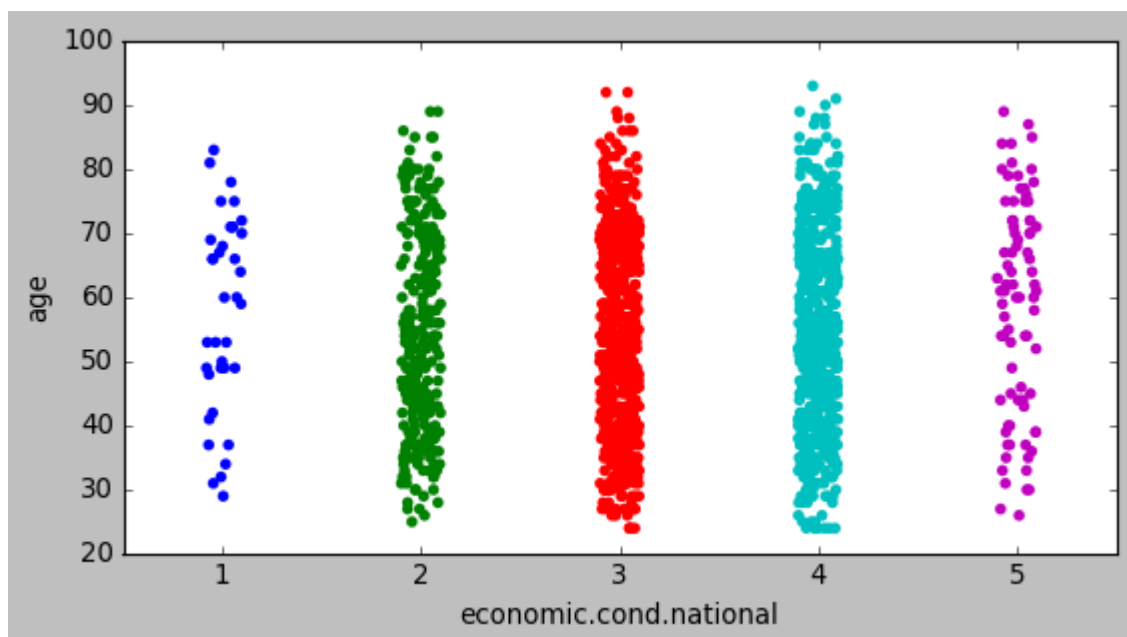
```
In [996]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["vote"],df['age'],jitter = True)
```

```
Out[996]: <AxesSubplot:xlabel='vote', ylabel='age'>
```




```
In [997]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["economic.cond.national"],df['age'],jitter = True)
```

```
Out[997]: <AxesSubplot:xlabel='economic.cond.national', ylabel='age'>
```



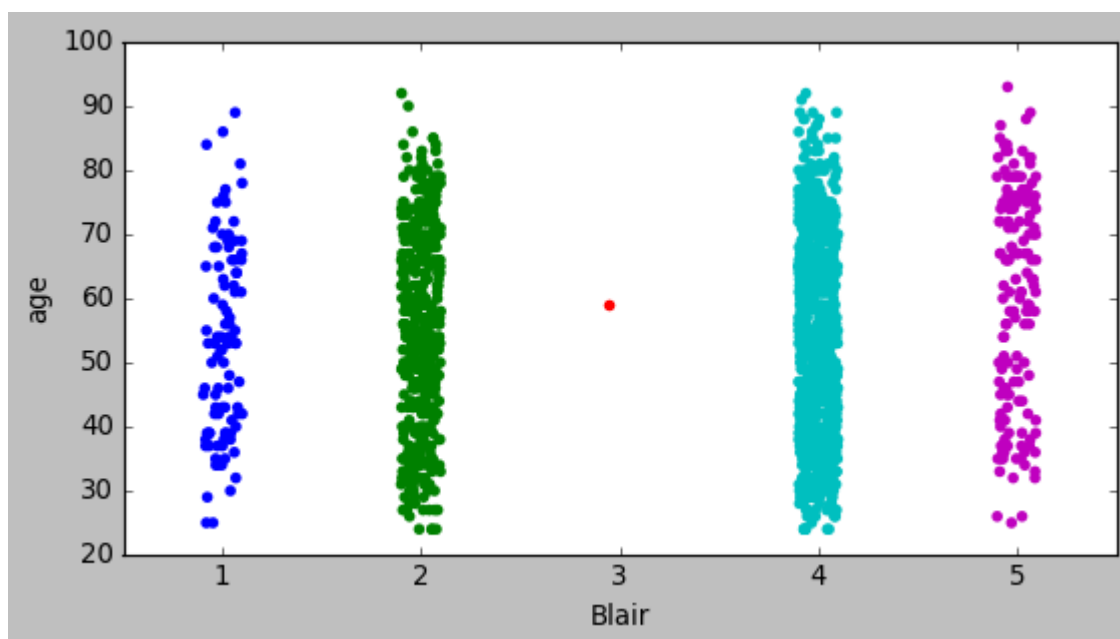
```
In [998]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["economic.cond.household"],df['age'],jitter = True)
```

Out[998]: <AxesSubplot:xlabel='economic.cond.household', ylabel='age'>



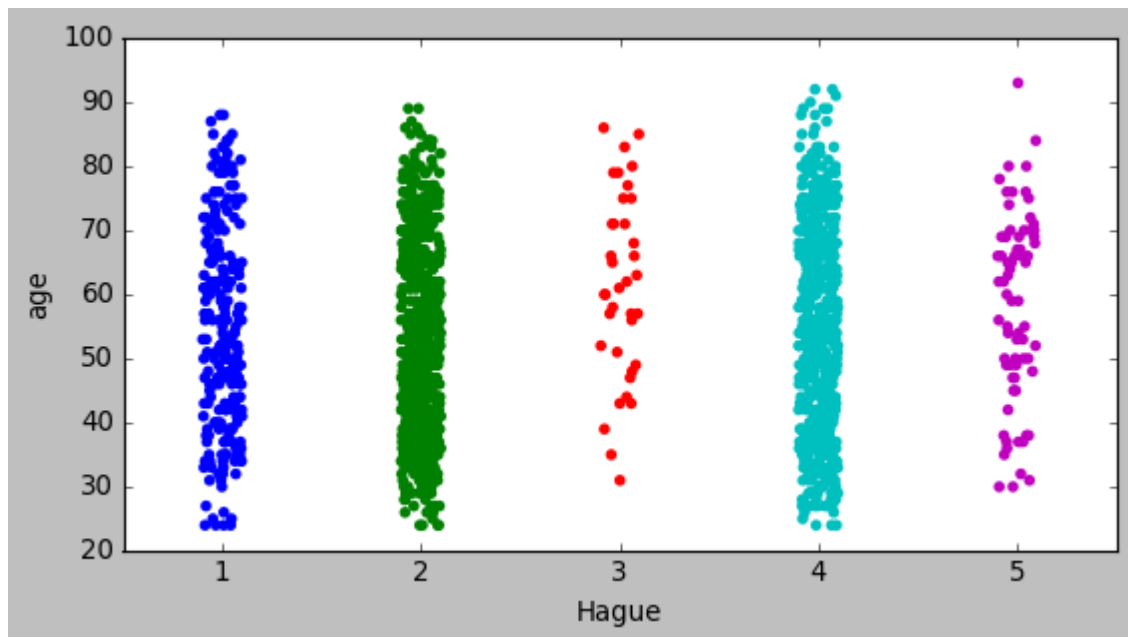
```
In [999]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["Blair"],df['age'],jitter = True)
```

Out[999]: <AxesSubplot:xlabel='Blair', ylabel='age'>




```
In [1000]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["Hague"],df['age'],jitter = True)
```

```
Out[1000]: <AxesSubplot:xlabel='Hague', ylabel='age'>
```




```
In [1001]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["Europe"],df['age'],jitter = True)
```

```
Out[1001]: <AxesSubplot:xlabel='Europe', ylabel='age'>
```




```
In [1002]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["political.knowledge"],df['age'],jitter = True)
```

Out[1002]: <AxesSubplot:xlabel='political.knowledge', ylabel='age'>



```
In [1003]: fig = plt.figure(figsize = (8,4))  
sns.stripplot(df["gender"],df['age'],jitter = True)
```

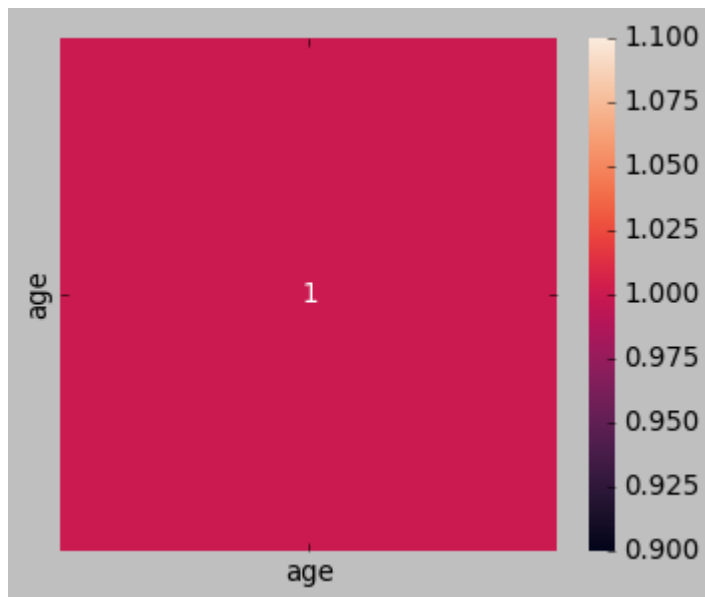
Out[1003]: <AxesSubplot:xlabel='gender', ylabel='age'>



In []:

Correlation Plot

```
In [1004]: plt.figure(figsize=(5,4))  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```



Since there is only one continuous variable, correlation cant be known for other features

```
In [1006]: df.corr()
```

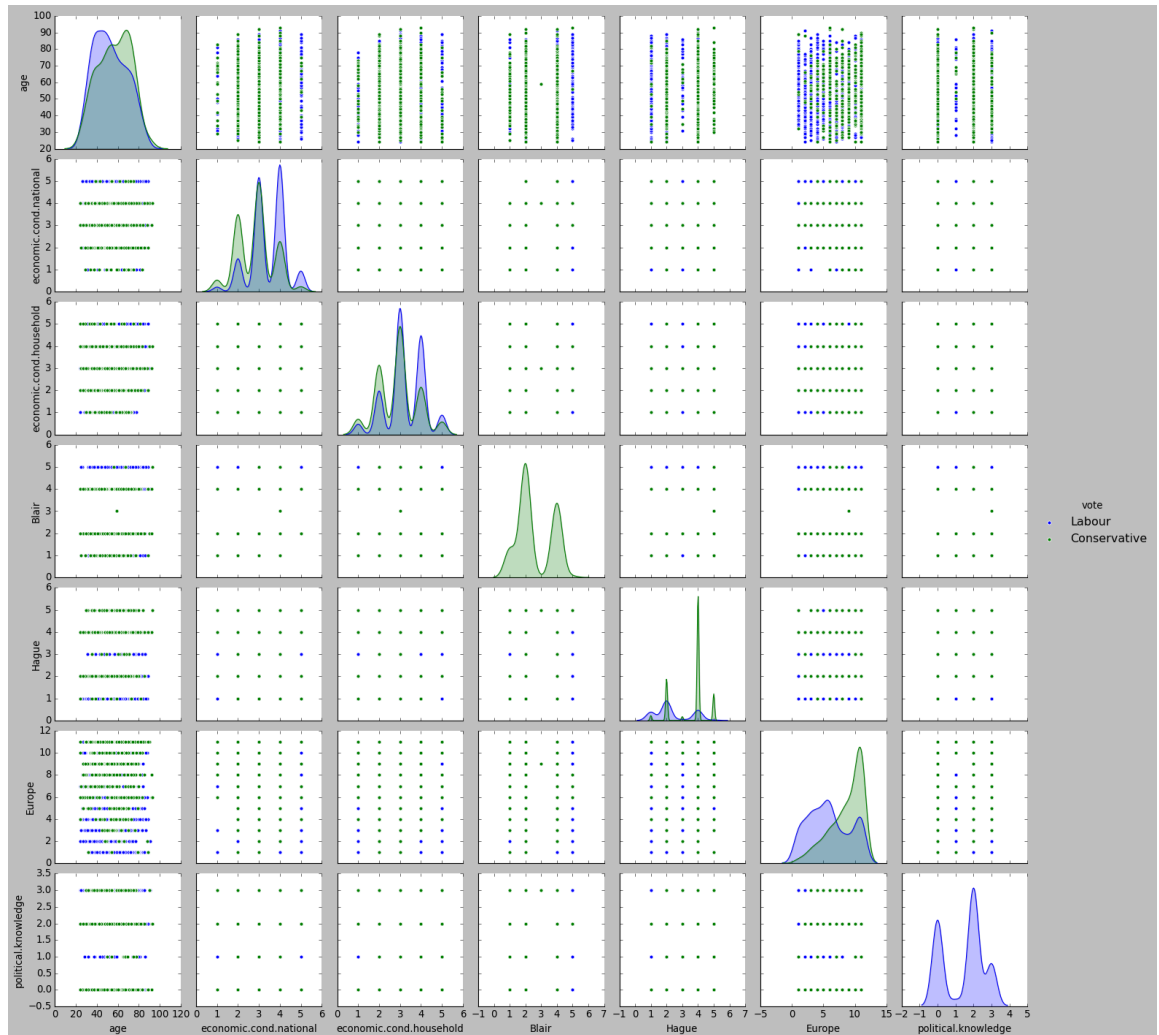
```
Out[1006]:
```

	age
age	1.0

```
In [1007]: plt.figure(figsize=(4,4))
sns.pairplot(df,hue = 'vote', diag_kind = 'kde')
```


Out[1007]: <seaborn.axisgrid.PairGrid at 0x168c995cfd0>

<Figure size 320x320 with 0 Axes>



Outlier Checks

```
In [1008]: # construct box plot for continuous variables  
plt.figure(figsize=(5,5))  
df.iloc[:,7].boxplot(vert=0)  
plt.show()
```




```
In [1009]: for feature in df.columns:
            if df[feature].dtype == 'object':
                print(feature)
                print(df[feature].value_counts())
                print('\n')
```

```
vote
Labour      1057
Conservative 460
Name: vote, dtype: int64
```

```
economic.cond.national
3      604
4      538
2      256
5       82
1       37
Name: economic.cond.national, dtype: int64
```

```
economic.cond.household
3      645
4      435
2      280
5       92
1       65
Name: economic.cond.household, dtype: int64
```

```
Blair
4      833
2      434
5      152
1       97
3        1
Name: Blair, dtype: int64
```

```
Hague
2      617
4      557
1      233
5       73
3       37
Name: Hague, dtype: int64
```

```
Europe
11     338
6     207
3     128
4     126
5     123
9     111
8     111
1     109
10    101
7      86
2      77
Name: Europe, dtype: int64
```

```
political.knowledge
2     776
0     454
3     249
```

```
1      38
Name: political.knowledge, dtype: int64
```

```
gender
female    808
male      709
Name: gender, dtype: int64
```

Data Preparation:

1. Encode the data (having string values) for Modelling. Is Scaling necessary here or not? Data Split: Split the data into train and test (70:30). **

Converting all objects to categorical codes



```
In [1010]: for feature in df.columns:
            if df[feature].dtype == 'object':
                print('\n')
                print('feature:', feature)
                print(pd.Categorical(df[feature].unique()))
                print(pd.Categorical(df[feature].unique()).codes)
                df[feature] = pd.Categorical(df[feature]).codes
```

```
feature: vote
['Labour', 'Conservative']
Categories (2, object): ['Conservative', 'Labour']
[1 0]
```

```
feature: economic.cond.national
[3, 4, 2, 1, 5]
Categories (5, int64): [1, 2, 3, 4, 5]
[2 3 1 0 4]
```

```
feature: economic.cond.household
[3, 4, 2, 1, 5]
Categories (5, int64): [1, 2, 3, 4, 5]
[2 3 1 0 4]
```

```
feature: Blair
[4, 5, 2, 1, 3]
Categories (5, int64): [1, 2, 3, 4, 5]
[3 4 1 0 2]
```

```
feature: Hague
[1, 4, 2, 5, 3]
Categories (5, int64): [1, 2, 3, 4, 5]
[0 3 1 4 2]
```

```
feature: Europe
[2, 5, 3, 4, 6, ..., 1, 7, 9, 10, 8]
Length: 11
Categories (11, int64): [1, 2, 3, 4, ..., 8, 9, 10, 11]
[ 1  4  2  3  5 10  0  6  8  9  7]
```

```
feature: political.knowledge
[2, 0, 3, 1]
Categories (4, int64): [0, 1, 2, 3]
[2 0 3 1]
```

```
feature: gender
['female', 'male']
Categories (2, object): ['female', 'male']
[0 1]
```

In [1011]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1517 entries, 0 to 1524
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   vote                                1517 non-null   int8
1   age                                1517 non-null   int64
2   economic.cond.national             1517 non-null   int8
3   economic.cond.household            1517 non-null   int8
4   Blair                              1517 non-null   int8
5   Hague                              1517 non-null   int8
6   Europe                             1517 non-null   int8
7   political.knowledge                 1517 non-null   int8
8   gender                             1517 non-null   int8
dtypes: int64(1), int8(8)
memory usage: 75.6 KB
```

In [1012]: `df.head()`

Out[1012]:

	vote	age	economic.cond.national	economic.cond.household	Blair	Hague	Europe	politic
0	1	43	2	2	3	0	1	
1	1	36	3	3	3	3	4	
2	1	35	3	3	4	1	2	
3	1	24	3	1	1	0	3	
4	1	41	1	1	0	0	5	

With Scaling

In [1013]: `cat1 = ['economic.cond.national', 'economic.cond.household', 'Blair', 'Hague', 'Europe', 'political.knowledge', 'gender']`

Scaling the variables as continuous variables have different weightage using min-max technique

In [1014]: `df = pd.get_dummies(df, columns=cat1, drop_first=True)`

In [1015]: df.head()

Out[1015]:

	vote	age	economic.cond.national_1	economic.cond.national_2	economic.cond.national_3
0	1	43	0	1	0
1	1	36	0	0	1
2	1	35	0	0	1
3	1	24	0	0	1
4	1	41	1	0	0

5 rows × 32 columns

In [1016]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1517 entries, 0 to 1524
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   vote                                1517 non-null   int8
1   age                                1517 non-null   int64
2   economic.cond.national_1          1517 non-null   uint8
3   economic.cond.national_2          1517 non-null   uint8
4   economic.cond.national_3          1517 non-null   uint8
5   economic.cond.national_4          1517 non-null   uint8
6   economic.cond.household_1          1517 non-null   uint8
7   economic.cond.household_2          1517 non-null   uint8
8   economic.cond.household_3          1517 non-null   uint8
9   economic.cond.household_4          1517 non-null   uint8
10  Blair_1                            1517 non-null   uint8
11  Blair_2                            1517 non-null   uint8
12  Blair_3                            1517 non-null   uint8
13  Blair_4                            1517 non-null   uint8
14  Hague_1                            1517 non-null   uint8
15  Hague_2                            1517 non-null   uint8
16  Hague_3                            1517 non-null   uint8
17  Hague_4                            1517 non-null   uint8
18  Europe_1                           1517 non-null   uint8
19  Europe_2                           1517 non-null   uint8
20  Europe_3                           1517 non-null   uint8
21  Europe_4                           1517 non-null   uint8
22  Europe_5                           1517 non-null   uint8
23  Europe_6                           1517 non-null   uint8
24  Europe_7                           1517 non-null   uint8
25  Europe_8                           1517 non-null   uint8
26  Europe_9                           1517 non-null   uint8
27  Europe_10                          1517 non-null   uint8
28  political.knowledge_1              1517 non-null   uint8
29  political.knowledge_2              1517 non-null   uint8
30  political.knowledge_3              1517 non-null   uint8
31  gender_1                          1517 non-null   uint8
dtypes: int64(1), int8(1), uint8(30)
memory usage: 109.6 KB
```

```
In [1026]: print(num)
```

```
['age']
```

```
In [1027]: df[num] = df[num].apply(lambda x:(x-x.min()) / (x.max()-x.min()))
```

```
In [1028]: ## Check if the variables have been scaled or not  
df.head()
```

```
Out[1028]:
```

	vote	age	economic.cond.national_1	economic.cond.national_2	economic.cond.national_3
0	1	0.275362	0	1	0
1	1	0.173913	0	0	0
2	1	0.159420	0	0	0
3	1	0.000000	0	0	0
4	1	0.246377	1	0	0

5 rows × 32 columns

Train-Test Split

```
In [1029]: df.columns
```

```
Out[1029]: Index(['vote', 'age', 'economic.cond.national_1', 'economic.cond.national_2',  
                'economic.cond.national_3', 'economic.cond.national_4',  
                'economic.cond.household_1', 'economic.cond.household_2',  
                'economic.cond.household_3', 'economic.cond.household_4', 'Blair_1',  
                'Blair_2', 'Blair_3', 'Blair_4', 'Hague_1', 'Hague_2', 'Hague_3',  
                'Hague_4', 'Europe_1', 'Europe_2', 'Europe_3', 'Europe_4', 'Europe_5',  
                'Europe_6', 'Europe_7', 'Europe_8', 'Europe_9', 'Europe_10',  
                'political.knowledge_1', 'political.knowledge_2',  
                'political.knowledge_3', 'gender_1'],  
              dtype='object')
```


```
In [1030]: # Copy all the predictor variables into X dataframe  
X = df.drop('vote', axis=1)  
  
# Copy target into the y dataframe.  
y = df['vote']
```

```
In [1031]: X.head()
```

```
Out[1031]:
```

	age	economic.cond.national_1	economic.cond.national_2	economic.cond.national_3	e
0	0.275362	0	1	0	
1	0.173913	0	0	1	
2	0.159420	0	0	1	
3	0.000000	0	0	1	
4	0.246377	1	0	0	

5 rows × 31 columns



```
In [1032]: y.head()
```

```
Out[1032]:
```

0	1
1	1
2	1
3	1
4	1

Name: vote, dtype: int8

```
In [1033]: # Split X and y into training and test set in 70:30 ratio
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3 , r
```


In [1034]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1517 entries, 0 to 1524
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   vote                                1517 non-null   int8
1   age                                1517 non-null   float64
2   economic.cond.national_1          1517 non-null   uint8
3   economic.cond.national_2          1517 non-null   uint8
4   economic.cond.national_3          1517 non-null   uint8
5   economic.cond.national_4          1517 non-null   uint8
6   economic.cond.household_1         1517 non-null   uint8
7   economic.cond.household_2         1517 non-null   uint8
8   economic.cond.household_3         1517 non-null   uint8
9   economic.cond.household_4         1517 non-null   uint8
10  Blair_1                           1517 non-null   uint8
11  Blair_2                           1517 non-null   uint8
12  Blair_3                           1517 non-null   uint8
13  Blair_4                           1517 non-null   uint8
14  Hague_1                           1517 non-null   uint8
15  Hague_2                           1517 non-null   uint8
16  Hague_3                           1517 non-null   uint8
17  Hague_4                           1517 non-null   uint8
18  Europe_1                          1517 non-null   uint8
19  Europe_2                          1517 non-null   uint8
20  Europe_3                          1517 non-null   uint8
21  Europe_4                          1517 non-null   uint8
22  Europe_5                          1517 non-null   uint8
23  Europe_6                          1517 non-null   uint8
24  Europe_7                          1517 non-null   uint8
25  Europe_8                          1517 non-null   uint8
26  Europe_9                          1517 non-null   uint8
27  Europe_10                         1517 non-null   uint8
28  political.knowledge_1              1517 non-null   uint8
29  political.knowledge_2              1517 non-null   uint8
30  political.knowledge_3              1517 non-null   uint8
31  gender_1                          1517 non-null   uint8
dtypes: float64(1), int8(1), uint8(30)
memory usage: 109.6 KB
```

Modelling

1. Apply Logistic Regression and LDA (linear discriminant analysis).

Logistic Regression

```
In [1035]: from sklearn.linear_model import LogisticRegression
from sklearn import metrics
from sklearn.metrics import roc_auc_score,roc_curve,classification_report,co
```

```
In [1036]: # Fit the Logistic Regression model
LR_model = LogisticRegression()
LR_model.fit(X_train, y_train)
```

```
Out[1036]: LogisticRegression()
```

Predicting on Training and Test dataset

```
In [1037]: ytrain_predict = LR_model.predict(X_train)
ytest_predict = LR_model.predict(X_test)
```

Getting the Predicted Classes and Probs

```
In [1038]: ytest_predict_proba=LR_model.predict_proba(X_test)
pd.DataFrame(ytest_predict_proba).head()
```

```
Out[1038]:
```

	0	1
0	0.641286	0.358714
1	0.231681	0.768319
2	0.021009	0.978991
3	0.908414	0.091586
4	0.107228	0.892772

Logistic Regression Model Evaluation

```
In [1039]: np.round(LR_model.coef_,decimals = 2)>0
```

```
Out[1039]: array([[False, False,  True,  True,  True, False,  True,  True, False,
        False, False,  True,  True, False, False, False, False,  True,
        True, False,  True, False, False, False, False, False, False,
        False, False, False,  True]])
```

```
In [1040]: from sklearn.feature_selection import RFE

predictor=X_train
selector = RFE(LR_model, n_features_to_select = 1)
selector = selector.fit(predictor,y_train)
selector.ranking_
```

```
Out[1040]: array([ 5, 14, 29,  7,  6, 19, 30, 23, 18, 13, 31,  4,  1, 20, 28,  3,  2,
        27, 22, 17, 21, 26, 12, 10,  8,  9, 11, 25, 15, 16, 24])
```

```
In [1041]: ## Performance Matrix on train data set
y_train_predict = LR_model.predict(X_train)
model_score = LR_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.8473138548539114

[[208 99]

[63 691]]

	precision	recall	f1-score	support
0	0.77	0.68	0.72	307
1	0.87	0.92	0.90	754
accuracy			0.85	1061
macro avg	0.82	0.80	0.81	1061
weighted avg	0.84	0.85	0.84	1061

```
In [1042]: ## Performance Matrix on test data set
y_test_predict = LR_model.predict(X_test)
model_score = LR_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8245614035087719

[[104 49]

[31 272]]

	precision	recall	f1-score	support
0	0.77	0.68	0.72	153
1	0.85	0.90	0.87	303
accuracy			0.82	456
macro avg	0.81	0.79	0.80	456
weighted avg	0.82	0.82	0.82	456

```
In [1068]: #the coefficients for each of the independent attributes

for idx, col_name in enumerate(X_train.columns):
    print("The coefficient for {} is {}".format(col_name, LR_model.coef_[0])
```

```
The coefficient for age is -1.1744662597904778
The coefficient for economic.cond.national_1 is -0.5176819122872669
The coefficient for economic.cond.national_2 is 0.056863771958550656
The coefficient for economic.cond.national_3 is 0.9439766626917371
The coefficient for economic.cond.national_4 is 1.0658858565395501
The coefficient for economic.cond.household_1 is -0.3636165042756915
The coefficient for economic.cond.household_2 is 0.0524779099905074
The coefficient for economic.cond.household_3 is 0.2851715158228671
The coefficient for economic.cond.household_4 is -0.4602130111030272
The coefficient for Blair_1 is -0.7831170915375294
The coefficient for Blair_2 is 0.0
The coefficient for Blair_3 is 0.6346235469331488
The coefficient for Blair_4 is 1.9405714934299025
The coefficient for Hague_1 is -0.4037435130688747
The coefficient for Hague_2 is -0.1127543404691797
The coefficient for Hague_3 is -1.9511780874524054
The coefficient for Hague_4 is -2.910718778649979
The coefficient for Europe_1 is 0.12530103447480292
The coefficient for Europe_2 is 0.21303170378589245
The coefficient for Europe_3 is -0.5400940404325377
The coefficient for Europe_4 is 0.29612321343401515
The coefficient for Europe_5 is -0.14341107825619337
The coefficient for Europe_6 is -0.5982110667160351
The coefficient for Europe_7 is -1.2983305031961327
The coefficient for Europe_8 is -1.4993756060970698
The coefficient for Europe_9 is -1.254106350023709
The coefficient for Europe_10 is -1.2406509400842118
The coefficient for political.knowledge_1 is -0.21622747500023098
The coefficient for political.knowledge_2 is -0.7309469728065002
The coefficient for political.knowledge_3 is -0.6399423403123045
The coefficient for gender_1 is 0.21054384787334637
```

```
In [1070]: # the intercept for the model

intercept = LR_model.intercept_[0]

print("The intercept for LR model is {}".format(intercept))
```

```
The intercept for LR model is 3.3258468981851084
```

```
In [1071]: # R square on testing data (coeff of determinant)
LR_model.score(X_test, y_test)
```

```
Out[1071]: 0.8245614035087719
```

```
In [1072]: # R square on training data
LR_model.score(X_train, y_train)
```

```
Out[1072]: 0.8473138548539114
```

```
In [1166]: #RMSE on Training data
predicted_train=LR_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1166]: 0.39075074554770667

```
In [1167]: #RMSE on Testing data
predicted_test=LR_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1167]: 0.4188539082916955

Discriminant Analysis

```
In [1043]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
LDA_model= LinearDiscriminantAnalysis()
LDA_model.fit(X_train, y_train)
```

Out[1043]: LinearDiscriminantAnalysis()

Predicting on Training and Test dataset

```
In [1044]: ytrain_predict = LDA_model.predict(X_train)
ytest_predict = LDA_model.predict(X_test)
```

Getting the Predicted Classes and Probs

```
In [1045]: ytest_predict_prob=LDA_model.predict_proba(X_test)
pd.DataFrame(ytest_predict_prob).head()
```

Out[1045]:

	0	1
0	0.676144	0.323856
1	0.184856	0.815144
2	0.012266	0.987734
3	0.953492	0.046508
4	0.079953	0.920047

LDA Model Evaluation

```
In [1046]: np.round(LDA_model.coef_,decimals = 2)>0
```

Out[1046]: array([[False, True, True, True, True, False, False, False, False,
 False, False, True, True, False, False, False, False, False,
 False, False, False, False, False, False, False, False, False,
 False, False, False, True]])

```
In [1047]: from sklearn.feature_selection import RFE

predictor=X_train
selector = RFE(LDA_model, n_features_to_select = 1)
selector = selector.fit(predictor,y_train)
selector.ranking_
```

```
Out[1047]: array([ 9, 30, 14, 11, 10, 19, 27, 29, 18, 15, 31,  4,  3, 24, 28,  2,  1,
                23, 21, 17, 22, 20, 16,  8,  5,  6,  7, 25, 12, 13, 26])
```

```
In [1048]: ## Performance Matrix on train data set
y_train_predict = LDA_model.predict(X_train)
model_score = LDA_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
0.8444863336475024
```

```
[[216  91]
```

```
 [ 74 680]]
```

	precision	recall	f1-score	support
0	0.74	0.70	0.72	307
1	0.88	0.90	0.89	754
accuracy			0.84	1061
macro avg	0.81	0.80	0.81	1061
weighted avg	0.84	0.84	0.84	1061

```
In [1050]: ## Performance Matrix on test data set
y_test_predict = LDA_model.predict(X_test)
model_score = LDA_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.8201754385964912
```

```
[[107  46]
```

```
 [ 36 267]]
```

	precision	recall	f1-score	support
0	0.75	0.70	0.72	153
1	0.85	0.88	0.87	303
accuracy			0.82	456
macro avg	0.80	0.79	0.79	456
weighted avg	0.82	0.82	0.82	456

```
In [1060]: #the coefficients for each of the independent attributes

for idx, col_name in enumerate(X_train.columns):
    print("The coefficient for {} is {}".format(col_name, LDA_model.coef_[0]
```

```
The coefficient for age is -1.6054817728067365
The coefficient for economic.cond.national_1 is 0.011057388537691917
The coefficient for economic.cond.national_2 is 0.8152787250715633
The coefficient for economic.cond.national_3 is 1.6444282399466414
The coefficient for economic.cond.national_4 is 1.8052982289827066
The coefficient for economic.cond.household_1 is -0.734342349289149
The coefficient for economic.cond.household_2 is -0.22969441775054084
The coefficient for economic.cond.household_3 is -0.08000247435303365
The coefficient for economic.cond.household_4 is -0.9087363797046396
The coefficient for Blair_1 is -0.6861893524294167
The coefficient for Blair_2 is -4.0483397333117293e-16
The coefficient for Blair_3 is 1.2397487082601484
The coefficient for Blair_4 is 2.075975766849963
The coefficient for Hague_1 is -0.44448043481909605
The coefficient for Hague_2 is -0.08488123913742465
The coefficient for Hague_3 is -2.4830826464084983
The coefficient for Hague_4 is -4.206117192482931
The coefficient for Europe_1 is -0.48194809336272615
The coefficient for Europe_2 is -0.555347309457755
The coefficient for Europe_3 is -1.288459692863717
The coefficient for Europe_4 is -0.4855053728879808
The coefficient for Europe_5 is -0.7387708641621507
The coefficient for Europe_6 is -1.25244993906857
The coefficient for Europe_7 is -2.2708058440884917
The coefficient for Europe_8 is -2.7585751578394517
The coefficient for Europe_9 is -2.31197928838158
The coefficient for Europe_10 is -2.1361281333340405
The coefficient for political.knowledge_1 is -0.38235599112698015
The coefficient for political.knowledge_2 is -1.068276354917667
The coefficient for political.knowledge_3 is -1.1207650745790063
The coefficient for gender_1 is 0.19959609113653495
```

```
In [1055]: # the intercept for the model

intercept = LDA_model.intercept_[0]

print("The intercept for LDA model is {}".format(intercept))
```

```
The intercept for LDA model is 4.241648712932388
```

```
In [1063]: # R square on testing data (coeff of determinant)
LDA_model.score(X_test, y_test)
```

```
Out[1063]: 0.8201754385964912
```

```
In [1064]: # R square on training data
LDA_model.score(X_train, y_train)
```

```
Out[1064]: 0.8444863336475024
```

```
In [1065]: #RMSE on Training data
predicted_train=LDA_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1065]: 0.39435221104045765

```
In [1067]: #RMSE on Testing data
predicted_test=LDA_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1067]: 0.4240572619393621

2. Apply KNN Model and Naïve Bayes Model. Interpret the results.

KNN Model

```
In [1075]: from sklearn.neighbors import KNeighborsClassifier

KNN_model=KNeighborsClassifier()
KNN_model.fit(X_train,y_train)
```

Out[1075]: KNeighborsClassifier()

Predicting on Training and Test dataset

```
In [1079]: ytrain_predict = KNN_model.predict(X_train)
ytest_predict = KNN_model.predict(X_test)
```

Getting the Predicted Classes and Probs

```
In [1080]: ytest_predict_prob=KNN_model.predict_proba(X_test)
pd.DataFrame(ytest_predict_prob).head()
```

```
Out[1080]:
```

	0	1
0	0.2	0.8
1	0.0	1.0
2	0.0	1.0
3	1.0	0.0
4	0.2	0.8

KNN Model Evaluation

```
In [1081]: ## Performance Matrix on train data set
y_train_predict = KNN_model.predict(X_train)
model_score = KNN_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.8501413760603205

[[212 95]

[64 690]]

	precision	recall	f1-score	support
0	0.77	0.69	0.73	307
1	0.88	0.92	0.90	754
accuracy			0.85	1061
macro avg	0.82	0.80	0.81	1061
weighted avg	0.85	0.85	0.85	1061

```
In [842]: ## Performance Matrix on test data set
y_test_predict = KNN_model.predict(X_test)
model_score = KNN_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.7828947368421053

[[91 62]

[37 266]]

	precision	recall	f1-score	support
0	0.71	0.59	0.65	153
1	0.81	0.88	0.84	303
accuracy			0.78	456
macro avg	0.76	0.74	0.75	456
weighted avg	0.78	0.78	0.78	456

```
In [1088]: # R square on testing data (coeff of determinant)
KNN_model.score(X_test, y_test)
```

Out[1088]: 0.7828947368421053

```
In [1089]: # R square on training data
KNN_model.score(X_train, y_train)
```

Out[1089]: 0.8501413760603205

```
In [1090]: #RMSE on Training data
predicted_train=KNN_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1090]: 0.38711577588581886

```
In [1091]: #RMSE on Testing data
predicted_test=KNN_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1091]: 0.46594555814804667

Naive Bayes Model

```
In [1082]: from sklearn.naive_bayes import GaussianNB
from sklearn import metrics
```

```
In [1083]: NB_model = GaussianNB()
NB_model.fit(X_train, y_train)
```

Out[1083]: GaussianNB()

Predicting on Training and Test dataset

```
In [1084]: ytrain_predict = NB_model.predict(X_train)
ytest_predict = NB_model.predict(X_test)
```

Getting the Predicted Classes and Probs

```
In [1085]: ytest_predict_prob=NB_model.predict_proba(X_test)
pd.DataFrame(ytest_predict_prob).head()
```

Out[1085]:

	0	1
0	9.954622e-01	0.004538
1	8.951301e-01	0.104870
2	9.487741e-32	1.000000
3	9.999868e-01	0.000013
4	1.345917e-09	1.000000

Naives Model Evaluation

```
In [1097]: ## Performance Matrix on train data set
y_train_predict = NB_model.predict(X_train)
model_score = NB_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.7492931196983977

[[248 59]

[207 547]]

	precision	recall	f1-score	support
0	0.55	0.81	0.65	307
1	0.90	0.73	0.80	754
accuracy			0.75	1061
macro avg	0.72	0.77	0.73	1061
weighted avg	0.80	0.75	0.76	1061

```
In [1098]: y_test.value_counts()
```

```
Out[1098]: 1    303
0    153
Name: vote, dtype: int64
```

```
In [1099]: ## Performance Matrix on test data set
y_test_predict = NB_model.predict(X_test)
model_score = NB_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.7346491228070176

[[120 33]

[88 215]]

	precision	recall	f1-score	support
0	0.58	0.78	0.66	153
1	0.87	0.71	0.78	303
accuracy			0.73	456
macro avg	0.72	0.75	0.72	456
weighted avg	0.77	0.73	0.74	456

```
In [1100]: #R square on testing data (coeff of determinant)
NB_model.score(X_test, y_test)
```

```
Out[1100]: 0.7346491228070176
```

```
In [1101]: # R square on training data
NB_model.score(X_train, y_train)
```

Out[1101]: 0.7492931196983977

```
In [1095]: #RMSE on Training data
predicted_train=NB_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1095]: 0.500706381327023

```
In [1096]: #RMSE on Testing data
predicted_test=NB_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1096]: 0.5151221963699317

3. Model Tuning, Bagging (Random Forest should be applied for Bagging) and Boosting.

Ada Boost

```
In [1102]: from sklearn.ensemble import AdaBoostClassifier

ADB_model = AdaBoostClassifier(n_estimators=100,random_state=1)
ADB_model.fit(X_train,y_train)
```

Out[1102]: AdaBoostClassifier(n_estimators=100, random_state=1)

```
In [1103]: ## Performance Matrix on train data set
y_train_predict = ADB_model.predict(X_train)
model_score = ADB_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.8473138548539114

[[211 96]

[66 688]]

	precision	recall	f1-score	support
0	0.76	0.69	0.72	307
1	0.88	0.91	0.89	754
accuracy			0.85	1061
macro avg	0.82	0.80	0.81	1061
weighted avg	0.84	0.85	0.84	1061

```
In [1104]: ## Performance Matrix on test data set
y_test_predict = ADB_model.predict(X_test)
model_score = ADB_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8135964912280702

[[100 53]

[32 271]]

	precision	recall	f1-score	support
0	0.76	0.65	0.70	153
1	0.84	0.89	0.86	303
accuracy			0.81	456
macro avg	0.80	0.77	0.78	456
weighted avg	0.81	0.81	0.81	456

```
In [1105]: #R square on testing data (coeff of determinant)
ADB_model.score(X_test, y_test)
```

Out[1105]: 0.8135964912280702

```
In [1106]: # R square on training data
ADB_model.score(X_train, y_train)
```

Out[1106]: 0.8473138548539114

```
In [1107]: #RMSE on Training data
predicted_train=ADB_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1107]: 0.39075074554770667

```
In [1108]: #RMSE on Testing data
predicted_test=ADB_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1108]: 0.43174472639735834

Gradient Boosting

```
In [1109]: from sklearn.ensemble import GradientBoostingClassifier
gbcl = GradientBoostingClassifier(random_state=1)
gbcl = gbcl.fit(X_train, y_train)
```

```
In [1110]: ## Performance Matrix on train data set
y_train_predict = gbcl.predict(X_train)
model_score = gbcl.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.884071630537229

[[227 80]

[43 711]]

	precision	recall	f1-score	support
0	0.84	0.74	0.79	307
1	0.90	0.94	0.92	754
accuracy			0.88	1061
macro avg	0.87	0.84	0.85	1061
weighted avg	0.88	0.88	0.88	1061

```
In [1111]: ## Performance Matrix on test data set
y_test_predict = gbcl.predict(X_test)
model_score = gbcl.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8223684210526315

[[101 52]

[29 274]]

	precision	recall	f1-score	support
0	0.78	0.66	0.71	153
1	0.84	0.90	0.87	303
accuracy			0.82	456
macro avg	0.81	0.78	0.79	456
weighted avg	0.82	0.82	0.82	456

```
In [1114]: #R square on testing data (coeff of determinant)
gbcl.score(X_test, y_test)
```

Out[1114]: 0.8223684210526315

```
In [1115]: # R square on training data
gbcl.score(X_train, y_train)
```

Out[1115]: 0.884071630537229

```
In [1116]: #RMSE on Training data
predicted_train=gbcl.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1116]: 0.3404825538302528

```
In [1117]: #RMSE on Testing data
predicted_test=gbcl.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1117]: 0.4214636152117623

Decision Tree

```
In [1118]: from sklearn import tree

DT_model= tree.DecisionTreeClassifier()
DT_model.fit(X_train, y_train)
```

Out[1118]: DecisionTreeClassifier()

```
In [1119]: ## Performance Matrix on train data set
y_train_predict = DT_model.predict(X_train)
model_score = DT_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
1.0
[[307  0]
 [ 0 754]]
```

		precision	recall	f1-score	support
	0	1.00	1.00	1.00	307
	1	1.00	1.00	1.00	754
accuracy				1.00	1061
macro avg		1.00	1.00	1.00	1061
weighted avg		1.00	1.00	1.00	1061

```
In [1120]: ## Performance Matrix on test data set
y_test_predict = DT_model.predict(X_test)
model_score = DT_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.756578947368421
[[ 93  60]
 [ 51 252]]
```

		precision	recall	f1-score	support
	0	0.65	0.61	0.63	153
	1	0.81	0.83	0.82	303
accuracy				0.76	456
macro avg		0.73	0.72	0.72	456
weighted avg		0.75	0.76	0.75	456

```
In [1121]: #R square on testing data (coeff of determinant)
DT_model.score(X_test, y_test)
```

Out[1121]: 0.756578947368421

```
In [1122]: # R square on training data
DT_model.score(X_train, y_train)
```

Out[1122]: 1.0

```
In [1123]: #RMSE on Training data
predicted_train=DT_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1123]: 0.0

```
In [1124]: #RMSE on Testing data
predicted_test=DT_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1124]: 0.4933771910329651

Random Forest

```
In [1125]: from sklearn.ensemble import RandomForestClassifier

RF_model=RandomForestClassifier(n_estimators=100,random_state=1)
RF_model.fit(X_train, y_train)
```

Out[1125]: RandomForestClassifier(random_state=1)

```
In [1126]: ## Performance Matrix on train data set
y_train_predict = RF_model.predict(X_train)
model_score = RF_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
1.0
[[307  0]
 [ 0 754]]
      precision    recall  f1-score   support

      0       1.00      1.00      1.00        307
      1       1.00      1.00      1.00        754

 accuracy          1.00          1.00          1.00        1061
 macro avg          1.00          1.00          1.00          1061
 weighted avg          1.00          1.00          1.00          1061
```



```
In [1127]: ## Performance Matrix on test data set
y_test_predict = RF_model.predict(X_test)
model_score = RF_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8026315789473685

```
[[ 93  60]
```

```
 [ 30 273]]
```

	precision	recall	f1-score	support
0	0.76	0.61	0.67	153
1	0.82	0.90	0.86	303
accuracy			0.80	456
macro avg	0.79	0.75	0.77	456
weighted avg	0.80	0.80	0.80	456

```
In [1128]: #R square on testing data (coeff of determinant)
RF_model.score(X_test, y_test)
```

Out[1128]: 0.8026315789473685

```
In [1129]: # R square on training data
RF_model.score(X_train, y_train)
```

Out[1129]: 1.0

```
In [1130]: #RMSE on Training data
predicted_train=RF_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1130]: 0.0

```
In [1131]: #RMSE on Testing data
predicted_test=RF_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1131]: 0.4442616583193193

Bagging

```
In [1137]: from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
cart = RandomForestClassifier()
Bagging_model=BaggingClassifier(base_estimator=cart,n_estimators=100,random_
Bagging_model.fit(X_train, y_train)
```

```
Out[1137]: BaggingClassifier(base_estimator=RandomForestClassifier(), n_estimators=10
0,
random_state=1)
```

```
In [1138]: ## Performance Matrix on train data set
y_train_predict = Bagging_model.predict(X_train)
model_score = Bagging_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.9679547596606974

[[277 30]

[4 750]]

	precision	recall	f1-score	support
0	0.99	0.90	0.94	307
1	0.96	0.99	0.98	754
accuracy			0.97	1061
macro avg	0.97	0.95	0.96	1061
weighted avg	0.97	0.97	0.97	1061

```
In [1139]: ## Performance Matrix on test data set
y_test_predict = Bagging_model.predict(X_test)
model_score = Bagging_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8179824561403509

[[97 56]

[27 276]]

	precision	recall	f1-score	support
0	0.78	0.63	0.70	153
1	0.83	0.91	0.87	303
accuracy			0.82	456
macro avg	0.81	0.77	0.78	456
weighted avg	0.81	0.82	0.81	456

SMOTE

```
In [1140]: from imblearn.over_sampling import SMOTE
```

SMOTE is only applied on the train data set

```
In [1141]: smt = SMOTE()
```

```
In [1143]: sm = SMOTE(random_state=2)
X_train_res, y_train_res = sm.fit_resample(X_train, y_train.ravel())
```

```
In [1144]: ## Let's check the shape after SMOTE
X_train_res.shape
```

```
Out[1144]: (1508, 31)
```

Naive Bayes with SMOTE

```
In [1145]: from sklearn.naive_bayes import GaussianNB
from sklearn import metrics
```

```
In [1146]: NB_SM_model = GaussianNB()
NB_SM_model.fit(X_train_res, y_train_res)
```

```
Out[1146]: GaussianNB()
```

```
In [1147]: ## Performance Matrix on train data set with SMOTE
y_train_predict = NB_SM_model.predict(X_train_res)
model_score = NB_SM_model.score(X_train_res, y_train_res)
print(model_score)
print(metrics.confusion_matrix(y_train_res, y_train_predict))
print(metrics.classification_report(y_train_res, y_train_predict))
```

```
0.7798408488063661
```

```
[[666  88]
```

```
 [244 510]]
```

	precision	recall	f1-score	support
0	0.73	0.88	0.80	754
1	0.85	0.68	0.75	754
accuracy			0.78	1508
macro avg	0.79	0.78	0.78	1508
weighted avg	0.79	0.78	0.78	1508

```
In [1148]: ## Performance Matrix on test data set
y_test_predict = NB_SM_model.predict(X_test)
model_score = NB_SM_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.6951754385964912
```

```
[[119  34]
```

```
 [105 198]]
```

	precision	recall	f1-score	support
0	0.53	0.78	0.63	153
1	0.85	0.65	0.74	303
accuracy			0.70	456
macro avg	0.69	0.72	0.69	456
weighted avg	0.75	0.70	0.70	456

KNN With SMOTE

In [1149]: `from sklearn.neighbors import KNeighborsClassifier`

```
KNN_SM_model=KNeighborsClassifier()
KNN_SM_model.fit(X_train_res,y_train_res)
```

Out[1149]: `KNeighborsClassifier()`

In [1150]: *## Performance Matrix on train data set*

```
y_train_predict = KNN_SM_model.predict(X_train_res)
model_score = KNN_SM_model.score(X_train_res, y_train_res)
print(model_score)
print(metrics.confusion_matrix(y_train_res, y_train_predict))
print(metrics.classification_report(y_train_res, y_train_predict))
```

0.8806366047745358

[[725 29]

[151 603]]

	precision	recall	f1-score	support
0	0.83	0.96	0.89	754
1	0.95	0.80	0.87	754
accuracy			0.88	1508
macro avg	0.89	0.88	0.88	1508
weighted avg	0.89	0.88	0.88	1508

In [1151]: *## Performance Matrix on test data set*

```
y_test_predict = KNN_SM_model.predict(X_test)
model_score = KNN_SM_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.743421052631579

[[116 37]

[80 223]]

	precision	recall	f1-score	support
0	0.59	0.76	0.66	153
1	0.86	0.74	0.79	303
accuracy			0.74	456
macro avg	0.72	0.75	0.73	456
weighted avg	0.77	0.74	0.75	456

Conclusion after SMOTE

Cross Validation on Naive Bayes Model

```
In [1152]: from sklearn.model_selection import cross_val_score
scores = cross_val_score(NB_SM_model, X_train_res, y_train_res, cv=10)
scores
```

```
Out[1152]: array([0.73509934, 0.78145695, 0.7615894 , 0.65562914, 0.80794702,
0.8013245 , 0.76821192, 0.82781457, 0.78      , 0.82      ])
```

```
In [1153]: scores = cross_val_score(NB_SM_model, X_test, y_test, cv=10)
scores
```

```
Out[1153]: array([0.56521739, 0.54347826, 0.67391304, 0.56521739, 0.54347826,
0.47826087, 0.57777778, 0.57777778, 0.6      , 0.53333333])
```

```
In [ ]: ## Cross Validation on KNN Model
```

```
In [1154]: from sklearn.model_selection import cross_val_score
scores = cross_val_score(KNN_SM_model, X_train_res, y_train_res, cv=10)
scores
```

```
Out[1154]: array([0.78807947, 0.8410596 , 0.84768212, 0.8013245 , 0.89403974,
0.8410596 , 0.87417219, 0.85430464, 0.81333333, 0.81333333])
```

```
In [1155]: scores = cross_val_score(KNN_SM_model, X_test, y_test, cv=10)
scores
```

```
Out[1155]: array([0.76086957, 0.73913043, 0.73913043, 0.73913043, 0.80434783,
0.80434783, 0.66666667, 0.75555556, 0.68888889, 0.73333333])
```

```
In [113]: ## After 10 fold cross validation, scores both on train and test data set re
## Hence our model is valid.
```

4. Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC_AUC score for each model. Final Model: Compare the models and write inference which model is best/optimized.

Logistic Regression

```
In [1156]: np.round(LR_model.coef_,decimals = 2)>0
```

```
Out[1156]: array([[False, False,  True,  True,  True, False,  True,  True, False,
        False, False,  True,  True, False, False, False, False,  True,
        True, False,  True, False, False, False, False, False, False,
        False, False, False,  True]])
```

```
In [1157]: from sklearn.feature_selection import RFE

predictor=X_train
selector = RFE(LR_model, n_features_to_select = 1)
selector = selector.fit(predictor,y_train)
selector.ranking_
```

```
Out[1157]: array([ 5, 14, 29,  7,  6, 19, 30, 23, 18, 13, 31,  4,  1, 20, 28,  3,  2,
        27, 22, 17, 21, 26, 12, 10,  8,  9, 11, 25, 15, 16, 24])
```

```
In [1158]: ## Performance Matrix on train data set
y_train_predict = LR_model.predict(X_train)
model_score = LR_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
0.8473138548539114
```

```
[[208  99]
```

```
 [ 63 691]]
```

	precision	recall	f1-score	support
0	0.77	0.68	0.72	307
1	0.87	0.92	0.90	754
accuracy			0.85	1061
macro avg	0.82	0.80	0.81	1061
weighted avg	0.84	0.85	0.84	1061

```
In [1159]: ## Performance Matrix on test data set
y_test_predict = LR_model.predict(X_test)
model_score = LR_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.8245614035087719
```

```
[[104  49]
```

```
 [ 31 272]]
```

	precision	recall	f1-score	support
0	0.77	0.68	0.72	153
1	0.85	0.90	0.87	303
accuracy			0.82	456
macro avg	0.81	0.79	0.80	456
weighted avg	0.82	0.82	0.82	456

```
In [1160]: #the coefficients for each of the independent attributes

for idx, col_name in enumerate(X_train.columns):
    print("The coefficient for {} is {}".format(col_name, LR_model.coef_[0])
```

```
The coefficient for age is -1.1744662597904778
The coefficient for economic.cond.national_1 is -0.5176819122872669
The coefficient for economic.cond.national_2 is 0.056863771958550656
The coefficient for economic.cond.national_3 is 0.9439766626917371
The coefficient for economic.cond.national_4 is 1.0658858565395501
The coefficient for economic.cond.household_1 is -0.3636165042756915
The coefficient for economic.cond.household_2 is 0.0524779099905074
The coefficient for economic.cond.household_3 is 0.2851715158228671
The coefficient for economic.cond.household_4 is -0.4602130111030272
The coefficient for Blair_1 is -0.7831170915375294
The coefficient for Blair_2 is 0.0
The coefficient for Blair_3 is 0.6346235469331488
The coefficient for Blair_4 is 1.9405714934299025
The coefficient for Hague_1 is -0.4037435130688747
The coefficient for Hague_2 is -0.1127543404691797
The coefficient for Hague_3 is -1.9511780874524054
The coefficient for Hague_4 is -2.910718778649979
The coefficient for Europe_1 is 0.12530103447480292
The coefficient for Europe_2 is 0.21303170378589245
The coefficient for Europe_3 is -0.5400940404325377
The coefficient for Europe_4 is 0.29612321343401515
The coefficient for Europe_5 is -0.14341107825619337
The coefficient for Europe_6 is -0.5982110667160351
The coefficient for Europe_7 is -1.2983305031961327
The coefficient for Europe_8 is -1.4993756060970698
The coefficient for Europe_9 is -1.254106350023709
The coefficient for Europe_10 is -1.2406509400842118
The coefficient for political.knowledge_1 is -0.21622747500023098
The coefficient for political.knowledge_2 is -0.7309469728065002
The coefficient for political.knowledge_3 is -0.6399423403123045
The coefficient for gender_1 is 0.21054384787334637
```

The sign of each coefficient indicates the direction of the relationship between a predictor variable and the response variable*

- Eg : For every 1 unit increase in Bair_3, vote increases by 0.635

For every 1 unit increase in Hague_3, vote decreases by 1.951

- Positive sign indicates that as the predictor variable increases the target variable also increases

- Negative sign indicates that as the predictor variable increases the target variable also decreases.

In []:

```
In [1161]: # the intercept for the model

intercept = LR_model.intercept_[0]

print("The intercept for LR model is {}".format(intercept))
```

The intercept for LR model is 3.3258468981851084

```
In [1162]: # R square on testing data (coeff of determinant)
LR_model.score(X_test, y_test)
```

Out[1162]: 0.8245614035087719

```
In [1163]: # R square on training data
LR_model.score(X_train, y_train)
```

Out[1163]: 0.8473138548539114

```
In [1170]: # RMSE on Training data
predicted_train=LR_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1170]: 0.39075074554770667


```
In [1168]: #RMSE on Testing data
predicted_test=LR_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1168]: 0.4188539082916955

AUC and ROC for training data

```
In [1169]: # predict probabilities
probs = LR_model.predict_proba(X_train)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(y_train, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr);
```


AUC: 0.903



AUC and ROC for Test data set

```
In [1171]: # predict probabilities
probs = LR_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.903



Linear Discriminant Analysis

```
In [1173]: np.round(LDA_model.coef_, decimals = 2)>0
```

```
Out[1173]: array([[False,  True,  True,  True,  True, False, False, False, False,
        False, False,  True,  True, False, False, False, False, False,
        False, False, False, False, False, False, False, False, False,
        False, False, False,  True]])
```

```
In [1175]: from sklearn.feature_selection import RFE

predictor=X_train
selector = RFE(LDA_model, n_features_to_select = 1)
selector = selector.fit(predictor,y_train)
selector.ranking_
```

```
Out[1175]: array([ 9, 30, 14, 11, 10, 19, 27, 29, 18, 15, 31,  4,  3, 24, 28,  2,  1,
                23, 21, 17, 22, 20, 16,  8,  5,  6,  7, 25, 12, 13, 26])
```

```
In [1177]: ## Performance Matrix on train data set
y_train_predict = LDA_model.predict(X_train)
model_score = LDA_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
0.8444863336475024
```

```
[[216  91]
```

```
 [ 74 680]]
```

	precision	recall	f1-score	support
0	0.74	0.70	0.72	307
1	0.88	0.90	0.89	754
accuracy			0.84	1061
macro avg	0.81	0.80	0.81	1061
weighted avg	0.84	0.84	0.84	1061

```
In [1178]: ## Performance Matrix on test data set
y_test_predict = LDA_model.predict(X_test)
model_score = LDA_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.8201754385964912
```

```
[[107  46]
```

```
 [ 36 267]]
```

	precision	recall	f1-score	support
0	0.75	0.70	0.72	153
1	0.85	0.88	0.87	303
accuracy			0.82	456
macro avg	0.80	0.79	0.79	456
weighted avg	0.82	0.82	0.82	456

In [1179]: *#the coefficients for each of the independent attributes*

```
for idx, col_name in enumerate(X_train.columns):
    print("The coefficient for {} is {}".format(col_name, LDA_model.coef_[0]
```

```
The coefficient for age is -1.6054817728067365
The coefficient for economic.cond.national_1 is 0.011057388537691917
The coefficient for economic.cond.national_2 is 0.8152787250715633
The coefficient for economic.cond.national_3 is 1.6444282399466414
The coefficient for economic.cond.national_4 is 1.8052982289827066
The coefficient for economic.cond.household_1 is -0.734342349289149
The coefficient for economic.cond.household_2 is -0.22969441775054084
The coefficient for economic.cond.household_3 is -0.08000247435303365
The coefficient for economic.cond.household_4 is -0.9087363797046396
The coefficient for Blair_1 is -0.6861893524294167
The coefficient for Blair_2 is -4.0483397333117293e-16
The coefficient for Blair_3 is 1.2397487082601484
The coefficient for Blair_4 is 2.075975766849963
The coefficient for Hague_1 is -0.44448043481909605
The coefficient for Hague_2 is -0.08488123913742465
The coefficient for Hague_3 is -2.4830826464084983
The coefficient for Hague_4 is -4.206117192482931
The coefficient for Europe_1 is -0.48194809336272615
The coefficient for Europe_2 is -0.555347309457755
The coefficient for Europe_3 is -1.288459692863717
The coefficient for Europe_4 is -0.4855053728879808
The coefficient for Europe_5 is -0.7387708641621507
The coefficient for Europe_6 is -1.25244993906857
The coefficient for Europe_7 is -2.2708058440884917
The coefficient for Europe_8 is -2.7585751578394517
The coefficient for Europe_9 is -2.31197928838158
The coefficient for Europe_10 is -2.1361281333340405
The coefficient for political.knowledge_1 is -0.38235599112698015
The coefficient for political.knowledge_2 is -1.068276354917667
The coefficient for political.knowledge_3 is -1.1207650745790063
The coefficient for gender_1 is 0.19959609113653495
```

In [1180]: *# the intercept for the model*

```
intercept = LDA_model.intercept_[0]

print("The intercept for LR model is {}".format(intercept))
```

The intercept for LR model is 4.241648712932388

In [1181]: *# R square on testing data (coeff of determinant)*

```
LDA_model.score(X_test, y_test)
```

Out[1181]: 0.8201754385964912

In [1182]: *# R square on training data*

```
LDA_model.score(X_train, y_train)
```

Out[1182]: 0.8444863336475024

```
In [1183]: # RMSE on Training data
predicted_train=LDA_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1183]: 0.39435221104045765


```
In [1184]: #RMSE on Testing data
predicted_test=LDA_model.fit(X_train, y_train).predict(X_test)
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1184]: 0.4240572619393621

AUC and ROC for Training Data

```
In [1185]: # predict probabilities
probs = LDA_model.predict_proba(X_train)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
auc = roc_auc_score(y_train, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(train_fpr, train_tpr);
```


AUC: 0.902



AUC and ROC for Test Data

```
In [1186]: # predict probabilities
probs = LDA_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.902



KNN Model

```
In [1188]: ## Performance Matrix on train data set
y_train_predict = KNN_model.predict(X_train)
model_score = KNN_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.8501413760603205

[[212 95]

[64 690]]

	precision	recall	f1-score	support
0	0.77	0.69	0.73	307
1	0.88	0.92	0.90	754
accuracy			0.85	1061
macro avg	0.82	0.80	0.81	1061
weighted avg	0.85	0.85	0.85	1061

```
In [1189]: ## Performance Matrix on test data set
y_test_predict = KNN_model.predict(X_test)
model_score = KNN_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.7828947368421053

[[91 62]

[37 266]]

	precision	recall	f1-score	support
0	0.71	0.59	0.65	153
1	0.81	0.88	0.84	303
accuracy			0.78	456
macro avg	0.76	0.74	0.75	456
weighted avg	0.78	0.78	0.78	456

```
In [1191]: # R square on testing data (coeff of determinant)
KNN_model.score(X_test, y_test)
```

Out[1191]: 0.7828947368421053

```
In [1192]: # R square on training data
KNN_model.score(X_train, y_train)
```

Out[1192]: 0.8501413760603205

```
In [1193]: # RMSE on Training data
predicted_train=KNN_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1193]: 0.38711577588581886


```
In [1194]: #RMSE on Testing data  
predicted_test=KNN_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1194]: 0.46594555814804667

AUC and ROC on the training data

```
In [1190]: # predict probabilities  
probs =KNN_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.924



AUC and ROC for Test data

```
In [1196]: # predict probabilities
probs = KNN_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.924



Naives Bayes Model

```
In [1198]: ## Performance Matrix on train data set
y_train_predict = NB_model.predict(X_train)
model_score = NB_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.7492931196983977

[[248 59]

[207 547]]

	precision	recall	f1-score	support
0	0.55	0.81	0.65	307
1	0.90	0.73	0.80	754
accuracy			0.75	1061
macro avg	0.72	0.77	0.73	1061
weighted avg	0.80	0.75	0.76	1061

```
In [1199]: ## Performance Matrix on test data set
y_test_predict = NB_model.predict(X_test)
model_score = NB_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.7346491228070176

[[120 33]

[88 215]]

	precision	recall	f1-score	support
0	0.58	0.78	0.66	153
1	0.87	0.71	0.78	303
accuracy			0.73	456
macro avg	0.72	0.75	0.72	456
weighted avg	0.77	0.73	0.74	456

```
In [1200]: # R square on testing data (coeff of determinant)
NB_model.score(X_test, y_test)
```

Out[1200]: 0.7346491228070176

```
In [1201]: # R square on training data
NB_model.score(X_train, y_train)
```

Out[1201]: 0.7492931196983977

```
In [1202]: # RMSE on Training data
predicted_train=NB_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1202]: 0.500706381327023


```
In [1203]: #RMSE on Testing data  
predicted_test=NB_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1203]: 0.5151221963699317

AUC and ROC for Training Data

```
In [1204]: # predict probabilities  
probs =NB_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.843



AUC and ROC for test data

```
In [1205]: # predict probabilities
probs = NB_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.843



Ada Boost

```
In [1207]: ## Performance Matrix on train data set
y_train_predict = ADB_model.predict(X_train)
model_score = ADB_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.8473138548539114

[[211 96]

[66 688]]

	precision	recall	f1-score	support
0	0.76	0.69	0.72	307
1	0.88	0.91	0.89	754
accuracy			0.85	1061
macro avg	0.82	0.80	0.81	1061
weighted avg	0.84	0.85	0.84	1061

```
In [1214]: ## Performance Matrix on test data set
y_test_predict = ADB_model.predict(X_test)
model_score = ADB_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8135964912280702

[[100 53]

[32 271]]

	precision	recall	f1-score	support
0	0.76	0.65	0.70	153
1	0.84	0.89	0.86	303
accuracy			0.81	456
macro avg	0.80	0.77	0.78	456
weighted avg	0.81	0.81	0.81	456

```
In [1228]: # R square on testing data (coeff of determinant)
ADB_model.score(X_test, y_test)
```

Out[1228]: 0.8135964912280702

```
In [1236]: # R square on training data
ADB_model.score(X_train, y_train)
```

Out[1236]: 0.8473138548539114

```
In [1244]: # RMSE on Training data
predicted_train=ADB_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1244]: 0.39075074554770667


```
In [1252]: #RMSE on Testing data  
predicted_test=ADB_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1252]: 0.43174472639735834

AUC and ROC for training data

```
In [1259]: # predict probabilities  
probs =ADB_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.912



AUC and ROC for test data

```
In [1266]: # predict probabilities
probs = ADB_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



Gradient Boost

```
In [1208]: ## Performance Matrix on train data set
y_train_predict = gbcl.predict(X_train)
model_score = gbcl.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.884071630537229

[[227 80]

[43 711]]

	precision	recall	f1-score	support
0	0.84	0.74	0.79	307
1	0.90	0.94	0.92	754
accuracy			0.88	1061
macro avg	0.87	0.84	0.85	1061
weighted avg	0.88	0.88	0.88	1061

```
In [1215]: ## Performance Matrix on test data set
y_test_predict = gbcl.predict(X_test)
model_score = gbcl.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8223684210526315

[[101 52]

[29 274]]

	precision	recall	f1-score	support
0	0.78	0.66	0.71	153
1	0.84	0.90	0.87	303
accuracy			0.82	456
macro avg	0.81	0.78	0.79	456
weighted avg	0.82	0.82	0.82	456

```
In [1227]: # R square on testing data (coeff of determinant)
gbcl.score(X_test, y_test)
```

Out[1227]: 0.8223684210526315

```
In [1235]: # R square on training data
gbcl.score(X_train, y_train)
```

Out[1235]: 0.884071630537229

```
In [1242]: # RMSE on Training data
predicted_train=gbcl.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1242]: 0.3404825538302528



```
In [1250]: #RMSE on Testing data  
predicted_test=gbcl.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1250]: 0.4214636152117623

AUC and ROC for training data

```
In [1258]: # predict probabilities  
probs =gbcl.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.945



AUC and ROC for test data

```
In [1265]: # predict probabilities
probs = gbc1.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



Decision Tree

In []:

In []:

```
In [1209]: ## Performance Matrix on train data set
y_train_predict = DT_model.predict(X_train)
model_score = DT_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
1.0
[[307  0]
 [ 0 754]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	307
1	1.00	1.00	1.00	754
accuracy			1.00	1061
macro avg	1.00	1.00	1.00	1061
weighted avg	1.00	1.00	1.00	1061

```
In [1216]: ## Performance Matrix on test data set
y_test_predict = DT_model.predict(X_test)
model_score = DT_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.756578947368421
[[ 94  59]
 [ 52 251]]
```

	precision	recall	f1-score	support
0	0.64	0.61	0.63	153
1	0.81	0.83	0.82	303
accuracy			0.76	456
macro avg	0.73	0.72	0.72	456
weighted avg	0.75	0.76	0.76	456

```
In [1226]: # R square on testing data (coeff of determinant)
DT_model.score(X_test, y_test)
```

Out[1226]: 0.756578947368421

```
In [1233]: # R square on training data
DT_model.score(X_train, y_train)
```

Out[1233]: 1.0

```
In [1241]: # RMSE on Training data
predicted_train=DT_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1241]: 0.0


```
In [1249]: #RMSE on Testing data  
predicted_test=DT_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1249]: 0.49559462778335206

AUC and ROC for training data

```
In [1257]: # predict probabilities  
probs =DT_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 1.000



AUC and ROC for test data

```
In [1264]: # predict probabilities
probs = DT_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



Random Forest

```
In [1210]: ## Performance Matrix on train data set
y_train_predict = RF_model.predict(X_train)
model_score = RF_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

```
1.0
[[307  0]
 [ 0 754]]
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	307
1	1.00	1.00	1.00	754
accuracy			1.00	1061
macro avg	1.00	1.00	1.00	1061
weighted avg	1.00	1.00	1.00	1061

```
In [1217]: ## Performance Matrix on test data set
y_test_predict = RF_model.predict(X_test)
model_score = RF_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

```
0.8026315789473685
[[ 93  60]
 [ 30 273]]
```

	precision	recall	f1-score	support
0	0.76	0.61	0.67	153
1	0.82	0.90	0.86	303
accuracy			0.80	456
macro avg	0.79	0.75	0.77	456
weighted avg	0.80	0.80	0.80	456

```
In [1225]: # R square on testing data (coeff of determinant)
RF_model.score(X_test, y_test)
```

Out[1225]: 0.8026315789473685

```
In [1232]: # R square on training data
RF_model.score(X_train, y_train)
```

Out[1232]: 1.0

```
In [1240]: # RMSE on Training data
predicted_train=RF_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1240]: 0.0


```
In [1248]: #RMSE on Testing data  
predicted_test=RF_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1248]: 0.4442616583193193

AUC and ROC for training data

```
In [1256]: # predict probabilities  
probs =RF_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 1.000



AUC and ROC for test data

```
In [1263]: # predict probabilities
probs = RF_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



Bagging with Random Forest


```
In [1211]: ## Performance Matrix on train data set
y_train_predict = Bagging_model.predict(X_train)
model_score = Bagging_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.9679547596606974

[[277 30]

[4 750]]

	precision	recall	f1-score	support
0	0.99	0.90	0.94	307
1	0.96	0.99	0.98	754
accuracy			0.97	1061
macro avg	0.97	0.95	0.96	1061
weighted avg	0.97	0.97	0.97	1061

```
In [1218]: ## Performance Matrix on test data set
y_test_predict = Bagging_model.predict(X_test)
model_score = Bagging_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.8179824561403509

[[97 56]

[27 276]]

	precision	recall	f1-score	support
0	0.78	0.63	0.70	153
1	0.83	0.91	0.87	303
accuracy			0.82	456
macro avg	0.81	0.77	0.78	456
weighted avg	0.81	0.82	0.81	456

```
In [1224]: # R square on testing data (coeff of determinant)
Bagging_model.score(X_test, y_test)
```

Out[1224]: 0.8179824561403509

```
In [1231]: # R square on training data
Bagging_model.score(X_train, y_train)
```

Out[1231]: 0.9679547596606974

```
In [1239]: # RMSE on Training data
predicted_train=Bagging_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1239]: 0.17901184413133828


```
In [1247]: #RMSE on Testing data  
predicted_test=Bagging_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1247]: 0.426635141379199

AUC and ROC for training data

```
In [1255]: # predict probabilities  
probs =Bagging_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.998



AUC and ROC for Test data

```
In [1262]: # predict probabilities
probs = Bagging_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



KNN with SMOTE

```
In [1212]: ## Performance Matrix on train data set
y_train_predict = KNN_SM_model.predict(X_train)
model_score = KNN_SM_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.8360037700282752

[[284 23]

[151 603]]

	precision	recall	f1-score	support
0	0.65	0.93	0.77	307
1	0.96	0.80	0.87	754
accuracy			0.84	1061
macro avg	0.81	0.86	0.82	1061
weighted avg	0.87	0.84	0.84	1061

```
In [1219]: ## Performance Matrix on test data set
y_test_predict = KNN_SM_model.predict(X_test)
model_score = KNN_SM_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.743421052631579

[[116 37]

[80 223]]

	precision	recall	f1-score	support
0	0.59	0.76	0.66	153
1	0.86	0.74	0.79	303
accuracy			0.74	456
macro avg	0.72	0.75	0.73	456
weighted avg	0.77	0.74	0.75	456

```
In [1223]: # R square on testing data (coeff of determinant)
KNN_SM_model.score(X_test, y_test)
```

Out[1223]: 0.743421052631579

```
In [1230]: # R square on training data
KNN_SM_model.score(X_train, y_train)
```

Out[1230]: 0.8360037700282752

```
In [1238]: # RMSE on Training data
predicted_train=KNN_SM_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1238]: 0.38711577588581886


```
In [1246]: #RMSE on Testing data  
predicted_test=KNN_SM_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1246]: 0.46594555814804667

AUC and ROC for Training data

```
In [1254]: # predict probabilities  
probs =KNN_SM_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.924



AUC and ROC for test data

```
In [1261]: # predict probabilities
probs = KNN_SM_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



NB with SMOTE

```
In [1213]: ## Performance Matrix on train data set
y_train_predict = NB_SM_model.predict(X_train)
model_score = NB_SM_model.score(X_train, y_train)
print(model_score)
print(metrics.confusion_matrix(y_train, y_train_predict))
print(metrics.classification_report(y_train, y_train_predict))
```

0.7115928369462771

[[245 62]

[244 510]]

	precision	recall	f1-score	support
0	0.50	0.80	0.62	307
1	0.89	0.68	0.77	754
accuracy			0.71	1061
macro avg	0.70	0.74	0.69	1061
weighted avg	0.78	0.71	0.72	1061

```
In [1267]: ## Performance Matrix on test data set
y_test_predict = NB_SM_model.predict(X_test)
model_score = KNN_SM_model.score(X_test, y_test)
print(model_score)
print(metrics.confusion_matrix(y_test, y_test_predict))
print(metrics.classification_report(y_test, y_test_predict))
```

0.7828947368421053

[[120 33]

[88 215]]

	precision	recall	f1-score	support
0	0.58	0.78	0.66	153
1	0.87	0.71	0.78	303
accuracy			0.73	456
macro avg	0.72	0.75	0.72	456
weighted avg	0.77	0.73	0.74	456

```
In [1222]: # R square on testing data (coeff of determinant)
NB_SM_model.score(X_test, y_test)
```

Out[1222]: 0.6951754385964912

```
In [1229]: # R square on training data
NB_SM_model.score(X_train, y_train)
```

Out[1229]: 0.7115928369462771

```
In [1237]: # RMSE on Training data
predicted_train=NB_SM_model.fit(X_train, y_train).predict(X_train)
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

Out[1237]: 0.500706381327023


```
In [1245]: #RMSE on Testing data  
predicted_test=NB_SM_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

Out[1245]: 0.5151221963699317

AUC and ROC for training data

```
In [1253]: # predict probabilities  
probs =NB_SM_model.predict_proba(X_train)  
# keep probabilities for the positive outcome only  
probs = probs[:, 1]  
# calculate AUC  
auc = roc_auc_score(y_train, probs)  
print('AUC: %.3f' % auc)  
# calculate roc curve  
train_fpr, train_tpr, train_thresholds = roc_curve(y_train, probs)  
plt.plot([0, 1], [0, 1], linestyle='--')  
# plot the roc curve for the model  
plt.plot(train_fpr, train_tpr);
```


AUC: 0.843



AUC and ROC for test data

```
In [1260]: # predict probabilities
probs = NB_SM_model.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
test_auc = roc_auc_score(y_test, probs)
print('AUC: %.3f' % auc)
# calculate roc curve
test_fpr, test_tpr, test_thresholds = roc_curve(y_test, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(test_fpr, test_tpr);
```

AUC: 0.912



In []: